

BigDataAnalysis BDA2

Helena Llorens Lluís (hllor282) Yi Yang (yiyang338)

May 2025

1 ASSIGNMENT 1

What are the lowest and highest temperatures measured each year for the period 1950- 2014. Provide the lists sorted in the descending order with respect to the maximum temperature. In this exercise you will use the temperature-readings.csv file.

- code

```
from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark.sql.functions import col, split, avg, max as spark_max, min as
↳ spark_min, countDistinct, count
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
↳ FloatType

# create a SparkContext which tells Spark how to access a cluster
spark = SparkSession.builder.appName("exercise 1").getOrCreate()
sc = spark.sparkContext

# create distributed datasets
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))
temps = lines.map(lambda x: (
    x[0],                # station number
    int(x[1][0:4]),       # year
    int(x[1][5:7]),       # month
    x[2],                # time
    float(x[3]),          # temperature
    x[4]                 # quality
))

# schema for the dataframe
schema = StructType([
    StructField("station_num", StringType()),
    StructField("year", IntegerType()),
    StructField("month", IntegerType()),
    StructField("time", StringType()),
```

```

        StructField("temp", FloatType()),
        StructField("quality", StringType()),
    ])
    df = spark.createDataFrame(temps, schema = schema)

    # filter the years
    filtered_df = df.where((col("year") >= 1950) & (col("year") <= 2014))

    # get max temperature per year
    max_temps = filtered_df.groupBy("year") \
        .agg(spark_max("temp").alias("temp"))

    # get the station number for the max temperatures
    max_temps = max_temps.join(filtered_df, ["year", "temp"]) \
        .select("year", "station_num", "temp") \
        .sort("temp", ascending = False)

    # get min temperature per year
    min_temps = filtered_df.groupBy("year") \
        .agg(spark_min("temp").alias("temp"))

    # get the station number for the min temperatures
    min_temps = min_temps.join(filtered_df, ["year", "temp"]) \
        .select("year", "station_num", "temp") \
        .sort("temp", ascending = False)

    # save result
    max_temps.write.csv("BDA/output/max_temps", header=True)
    min_temps.write.csv("BDA/output/min_temps", header=True)

```

- Maximum (*partial results*)

```

[
  year,station_num,temp
  1975,86200,36.1
  1992,63600,35.4
  1994,117160,34.7
  2014,96560,34.4
  2010,75250,34.4
  1989,63050,33.9
  1982,94050,33.8
  1968,137100,33.7
  1966,151640,33.5
  1983,98210,33.3
  2002,78290,33.3
  2002,78290,33.3
  1970,103080,33.2
  1986,76470,33.2
  2000,62400,33.0
  1956,145340,33.0
  1959,65160,32.8
  1991,137040,32.7
  2006,75240,32.7
  1988,102540,32.6
]

```

- Minimum (*partial results*)

```
[
  year,station_num,temp
1990,166870,-35.0
1990,147270,-35.0
1952,192830,-35.5
1974,179950,-35.6
1974,166870,-35.6
1954,113410,-36.0
1992,179960,-36.1
1975,157860,-37.0
1972,167860,-37.5
1995,182910,-37.6
2000,169860,-37.6
1957,159970,-37.8
1983,191900,-38.2
1989,166870,-38.2
1953,183760,-38.4
2009,179960,-38.5
1993,191900,-39.0
1984,123480,-39.2
1984,191900,-39.2
2008,179960,-39.3
]
```

2 ASSIGNMENT 2

Count the number of readings for each month in the period of 1950-2014 which are higher than 10 degrees.Repeat the exercise, this time taking only distinct readings from each station. That is, if a station reported a reading above 10 degrees in some month, then it appears only once in the count for that month. In this exercise, you will use the temperature-readings.csv file.

- code

```
from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark.sql.functions import col, split, avg, max as spark_max, min as
↳ spark_min, countDistinct, count
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
↳ FloatType

# create a SparkContext which tells Spark how to access a cluster
spark = SparkSession.builder.appName("exercise 2").getOrCreate()
sc = spark.sparkContext

# create distributed datasets
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))
```

```

temps = lines.map(lambda x: (
    x[0],          # station number
    int(x[1][0:4]), # year
    int(x[1][5:7]), # month
    x[2],          # time
    float(x[3]),   # temperature
    x[4]           # quality
))

# schema for the dataframe
schema = StructType([
    StructField("station_num", StringType()),
    StructField("year", IntegerType()),
    StructField("month", IntegerType()),
    StructField("time", StringType()),
    StructField("temp", FloatType()),
    StructField("quality", StringType()),
])
df = spark.createDataFrame(temps, schema = schema)

# filter years and temperature over 10°C
filtered_df = df.where((col("year") >= 1950) & (col("year") <= 2014) &
    ↪ (col("temp") > 10))

# count
count_temp = filtered_df.groupBy("year", "month") \
    .count().alias("count")

count_temp = count_temp.sort("count", ascending = False)

count_temp.write.csv("BDA/output/count")

# count distinct
count_temp = filtered_df.groupBy("year", "month") \
    .agg(countDistinct("station_num").alias("count_dist"))

count_temp = count_temp.sort("count_dist", ascending = False)

count_temp.write.csv("BDA/output/count_distinct")

```

-
- Count of readings(temperature higher than 10 degrees) per month (*partial results*)
-

```

[
  year,month,value
  2014,7,147681
  2011,7,146656
  2010,7,143419
  2012,7,137477
  2013,7,133657
  2009,7,133008
  2011,8,132734
  2009,8,128349
  2013,8,128235
  2003,7,128133

```

```

2002,7,127956
2006,8,127622
2008,7,126973
2002,8,126073
2005,7,125294
2011,6,125193
2012,8,125037
2006,7,124794
2010,8,124417
2014,8,124045
1997,7,123496
]

```

- Count of distinct readings(temperature higher than 10 degree) per month:
(*partial results*)

```

[
  year,month,value
1972,10,378
1973,5,377
1973,6,377
1973,9,376
1972,8,376
1972,6,375
1972,5,375
1971,8,375
1972,9,375
1971,6,374
1971,9,374
1972,7,374
1971,5,373
1973,8,373
1974,8,372
1974,6,372
1974,9,370
1970,8,370
1973,7,370
1974,5,370
1971,7,370
]

```

3 ASSIGNMENT 3

Find the average monthly temperature for each available station in Sweden. Your result should include average temperature for each station for each month in the period of 1960- 2014. Bear in mind that not every station has the readings for each month in this timeframe. In this exercise you will use the temperature-readings.csv file.

- code

```

from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark.sql.functions import col, split, avg, max as spark_max, min as
↳ spark_min, countDistinct, count
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
↳ FloatType

# create a SparkContext which tells Spark how to access a cluster
spark = SparkSession.builder.appName("exercise 3").getOrCreate()
sc = spark.sparkContext

# create distributed datasets
# This path is to the file on hdfs
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
lines = temperature_file.map(lambda line: line.split(";"))
temps = lines.map(lambda x: (
    x[0],                # station number
    int(x[1][0:4]),       # year
    int(x[1][5:7]),       # month
    x[2],                 # time
    float(x[3]),          # temperature
    x[4],                 # quality
))

# schema for the dataframe
schema = StructType([
    StructField("station_num", StringType()),
    StructField("year", IntegerType()),
    StructField("month", IntegerType()),
    StructField("time", StringType()),
    StructField("temp", FloatType()),
    StructField("quality", StringType()),
])
df = spark.createDataFrame(temps, schema = schema)

# filter years
filtered_df = df.where((col("year") >= 1960) & (col("year") <= 2014))

# average monthly temperature per station
avg_temp = filtered_df.groupBy("station_num", "year", "month") \
    .agg(avg("temp").alias("avg_temp"))

avg_temp = avg_temp.sort("avg_temp", ascending = False)

# save result
avg_temp.write.csv("BDA/output/avg_temp")

```

- The average monthly temperature for each available station. (*partial results*)

```
[
    station,year,month,avgMonthlyTemperature

```

```

96000,2014,7,26.299999237060547
65450,1994,7,23.654838664557346
95160,1994,7,23.50537642612252
75120,1994,7,23.268817204301076
105260,1994,7,23.143820259008514
85280,1994,7,23.108602175148583
54550,1983,8,23.0
54550,1975,8,22.9625
96550,1994,7,22.957894760265685
96000,1994,7,22.931182820309875
106070,1994,7,22.822580665670415
173960,1972,7,22.776666831970214
54300,1994,7,22.760215082476215
85210,1994,7,22.755913970290973
65450,2006,7,22.740860190442813
75120,2006,7,22.73010758430727
103080,1994,7,22.708602125926685
92100,1994,7,22.69892466965542
94180,1994,7,22.681720487533077
83230,1994,7,22.577419393806046
]

```

4 ASSIGNMENT 4

Provide a list of stations with their associated maximum measured temperatures and maximum measured daily precipitation. Show only those stations where the maximum temperature is between 25 and 30 degrees and maximum daily precipitation is between 100 mm and 200mm. In this exercise you will use the temperature-readings.csv and precipitation-readings.csv files.

- code
-

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import col, max as spark_max, sum as spark_sum
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
↳ FloatType

# Create a SparkSession
spark = SparkSession.builder.appName("exercise 4").getOrCreate()
sc = spark.sparkContext

# Read and parse temperature dataset
temperature_file = sc.textFile("BDA/input/temperature-readings.csv")
temp_lines = temperature_file.map(lambda x: x.split(";")).map(lambda x: (
    x[0],                # station_num
    x[1],                # date
    float(x[3])          # temperature
))

temp_schema = StructType([
    StructField("station_num", StringType()),

```

```

    StructField("date", StringType()),
    StructField("temp", FloatType()),
  ])

df_temp = spark.createDataFrame(temp_lines, schema=temp_schema)

# Get maximum temperature per station
max_temp_per_station = df_temp.groupBy("station_num") \
    .agg(spark_max("temp").alias("max_temp")) \
    .filter((col("max_temp") >= 25) & (col("max_temp")
    ↪ <= 30))

# Read and parse precipitation dataset
precipitation_file = sc.textFile("BDA/input/precipitation-readings.csv")
prec_lines = precipitation_file.map(lambda x: x.split(";")).map(lambda x: (
    x[0],                # station_num
    x[1],                # date
    float(x[3])          # precipitation
))

prec_schema = StructType([
    StructField("station_num", StringType()),
    StructField("date", StringType()),
    StructField("prec", FloatType()),
])

df_prec = spark.createDataFrame(prec_lines, schema=prec_schema)

# Sum daily precipitation per station
daily_prec = df_prec.groupBy("station_num", "date") \
    .agg(spark_sum("prec").alias("daily_prec"))

# Max daily precipitation per station
max_prec_per_station = daily_prec.groupBy("station_num") \
    .agg(spark_max("daily_prec").alias("max_prec"))
    ↪ \
    .filter((col("max_prec") >= 100) &
    ↪ (col("max_prec") <= 200))

# Join filtered temp and prec on station number
final_result = max_temp_per_station.join(max_prec_per_station, on="station_num",
    ↪ how="inner") \
    .select("station_num", "max_temp", "max_prec")
    ↪ \
    .orderBy("station_num", ascending=False)

# Save result
final_result.write.csv("BDA/output/exercise4", header=True)

```

No results were returned for the query filtering stations with a maximum temperature between 25°C and 30°C and a maximum daily precipitation between 100 mm and 200 mm. This suggests that no stations in the dataset satisfy both criteria.

5 ASSIGNMENT 5

Calculate the average monthly precipitation for the Östergötland region (list of stations is provided in a separate file) for the period 1993-2016. In order to do this, you will first need to calculate the total monthly precipitation for each station before calculating the monthly average (by averaging over stations). In this exercise you will use the precipitation-readings.csv and stations-Ostergotland.csv files. HINT (not for the SparkSQL lab): Avoid using joins here! stations-Ostergotland.csv is small and if distributed will cause a number of unnecessary shuffles when joined with precipitationRDD. If you distribute precipitation-readings.csv then either repartition your stations RDD to 1 partition or make use of the collect function to acquire a python list and broadcast function to broadcast the list to all nodes.

- code

```
from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark.sql.functions import col, split, avg, max as spark_max, min as
↳ spark_min, countDistinct, count, sum
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
↳ FloatType

# create a SparkContext which tells Spark how to access a cluster
spark = SparkSession.builder.appName("exercise 5").getOrCreate()
sc = spark.sparkContext

# create distributed datasets
# This path is to the file on hdfs

# precipitation file
precipitation_file = sc.textFile("BDA/input/precipitation-readings.csv")
prec_lines = precipitation_file.map(lambda line: line.split(";"))

prec_lines = prec_lines.map(lambda x: (
    x[0],                # stat_num
    int(x[1][0:4]),      # year
    int(x[1][5:7]),      # month
    x[2],                # time
    float(x[3]),         # precipitation (mm)
    x[4],                # quality
))

# precipitation dataset schema
prec_schema = StructType([
    StructField("station_num", StringType()),
    StructField("year", IntegerType()),
    StructField("month", IntegerType()),
    StructField("time", StringType()),
    StructField("prec", FloatType()),
    StructField("quality", StringType()),
])
```

```

df_prec = spark.createDataFrame(prec_lines, schema = prec_schema)

# stations Ostergotland dataset
oster_file = sc.textFile("BDA/input/stations-Ostergotland.csv")
oster_lines = oster_file.map(lambda line: line.split(";"))

oster_lines = oster_lines.map(lambda x: (
    x[0],                # stat_num
    x[1]                 # stat_name
))

# stations Ostergotland dataset schema
oster_schema = StructType([
    StructField("station_num", StringType()),
    StructField("station_name", StringType()),
])
df_oster = spark.createDataFrame(oster_lines, schema = oster_schema)

# filter years
filtered_df = df_prec.where((col("year") >= 1993) & (col("year") <= 2016))

# get only Ostergotland stations
stat_oster = filtered_df.join(df_oster, ["station_num"])

# average precipitation per station
avg_prec = stat_oster.groupBy("station_num", "year", "month") \
    .agg(sum("prec").alias("avg_prec"))

# average precipitation per month
avg_prec_month = avg_prec.groupBy("year", "month") \
    .agg(avg("avg_prec").alias("avg_prec_month")) \
    .orderBy(["year", "month"], ascending = False)

# save results
avg_prec_month.write.csv("BDA/output/avg_prec")

```

-
- Monthly precipitation average for stations in Ostergotland. (*partial results*)

```

[
  year,month,avgMonthlyPrecipitation
  2016,7,0.0
  2016,6,47.662500091828406
  2016,5,29.250000204890966
  2016,4,26.900000237859786
  2016,3,19.96250029373914
  2016,2,21.56250028219074
  2016,1,22.325000344775617
  2015,12,28.92500019259751
  2015,11,63.887500293552876
  2015,10,2.2625000346451998
  2015,9,101.30000030715019
  2015,8,26.9875001097098
  2015,7,119.09999992512167

```

2015,6,78.66250023245811
2015,5,93.22500019986182
2015,4,15.337500078603625
2015,3,42.612500292249024
2015,2,24.825000358745456
2015,1,59.11250052880496
2014,12,35.462500284425914
2014,11,52.42500038724393

]
